


MultiThreading Plugin

SOFA Training days 2013
Montpellier

Federico Spadoni, Hervé Delingette

Asclepios



Add multithreading to SOFA

Why multi-thread ?

- Speed-up computation (in addition to GPU)
- Keep interactive visualisation despite slow computation

Multithreading design:


- * Task and Scheduler (*KAAPI, intel TBB*)

Problems:

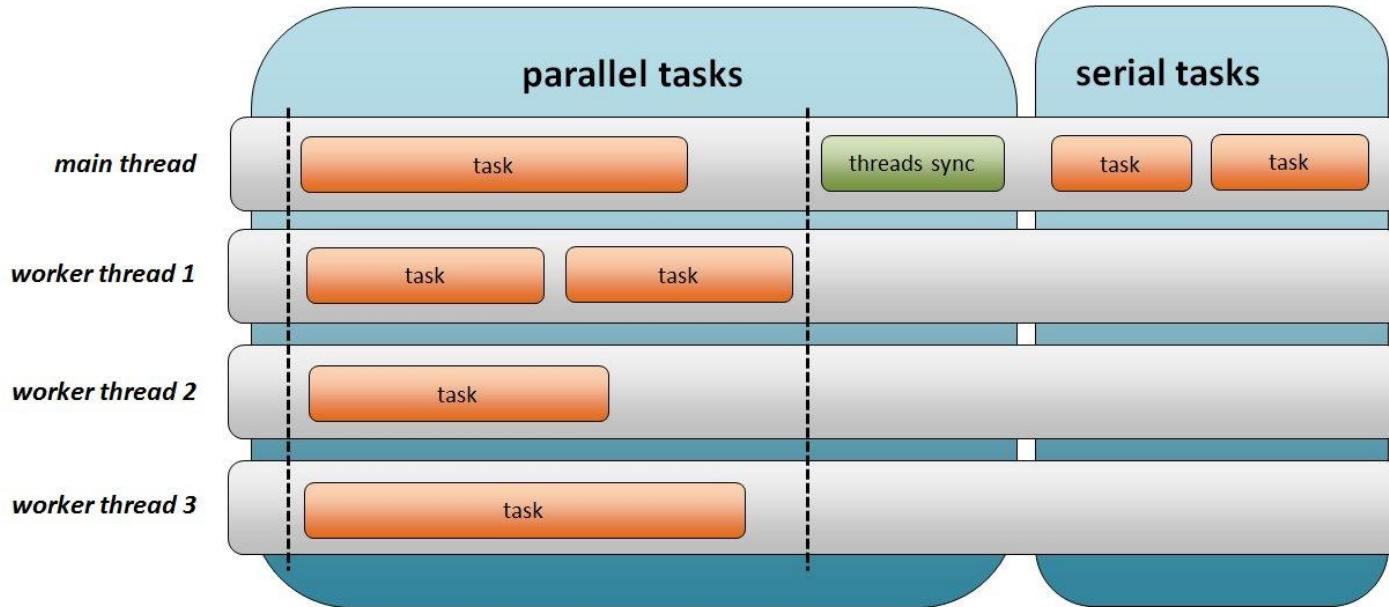
- ? Shared Data Synchronization

Constraint:


- ! Do not change Sofa core code
- ! Time



Task and Scheduler Overview




- the computation work load must be functionally decomposed into tasks.
- the scheduler maps the tasks onto physical threads




Task

- An independent functional block that can be executed in parallel
- Faster than a thread to create and destroy for the operating system
- It can itself create new tasks
- do not have data dependencies between other tasks executed concurrently



Task Scheduler (boost::thread)

- creates $n-1$ worker threads on a system with n logical cores, and takes care of their synchronization (idling, running).
- each worker thread manages its own queue of tasks.
 - every time one task finishes the worker thread pops the next one from the top of its queue
 - when a task is created it is pushed directly on to the top of the queue
- the scheduler steals the bottom half of the queue of a busy thread and gives it to a starving thread



Task implementation

```
#include "Task.h"


class MyTask : public Task
{
public:

    MyTask( Task::Status* status )
        : Task(status)
        {}

    virtual bool run(WorkerThread* )
    {
        // do something
    }

};
```


- Override the *run()* function
- Pass a *task::Status* pointer to the constructor



Task creation and queuing

```
{  
    Task::Status status;  
  
    MyTask* task = new MyTask( &status );  
  
    WorkerThread* curThread = WorkerThread::getCurrent();  
    curThread->addTask( MyTask );  
  
    ..  
    ..  
}
```


- The *addTask()* push *MyTask* in the queue of the current thread
- the current thread keeps executing the code and it can create and queue other tasks.
- Take care of the *MyTask* destruction



Task synchronization (blocking)

```
{  
    Task::Status status;  
  
    MyTask* task = new MyTask( &status );  
  
    WorkerThread* curThread = WorkerThread::getCurrent();  
  
    curThread->addTask( MyTask );  
  
    curThread->workUntilDone(&status);  
}
```

- *workUntilDone()* stops the execution until all the tasks with the same *status* completed
- the current thread starts to execute tasks from its queue or steals from other thread queue.



Task synchronization (not blocking 1)

```
{
    Task::Status status;


    MyTask* task = new MyTask(&status);

    WorkerThread* curThread = WorkerThread::getCurrent();

    curThread->addTask( MyTask );

    while ( status.IsBusy() )
    {
        // do something while tasks with
        // the same status flag complete
    }
}
```

- It executes some code inside the *while* loop concurrently with the tasks already queued until all the tasks with the same *status* completed (*status.IsBusy() == false*)



Task synchronization (not blocking 2)

```
{
    Task::Status status;

    while ( ;; )
    {
        if ( !status.IsBusy() )
        {

            MyTask* task = new MyTask(&status);

            WorkerThread* curThread = WorkerThread::getCurrent();


            curThread->addTask( MyTask );

        }

        // do something

    }
}
```

- The *if* statement block is executed when *MyTask* is not running
- The *while* loop queues continuously *MyTask* as soon as it completed



Creating tasks

- Component level:

running two or more components in parallel

AnimationLoopParallel runs all the *BaseAnimationLoop::step()* function found in the child nodes in parallel

- in functions:


Parallelizing some floating point computation inside a *for loop*

BeamLinearMapping_mt in *apply()* *applyJ()* *applyJT()*




AnimationLoopParallel and *DataExchange*

(component level parallelism)




```
<node "root">  
  <AnimationLoopParallel>
```

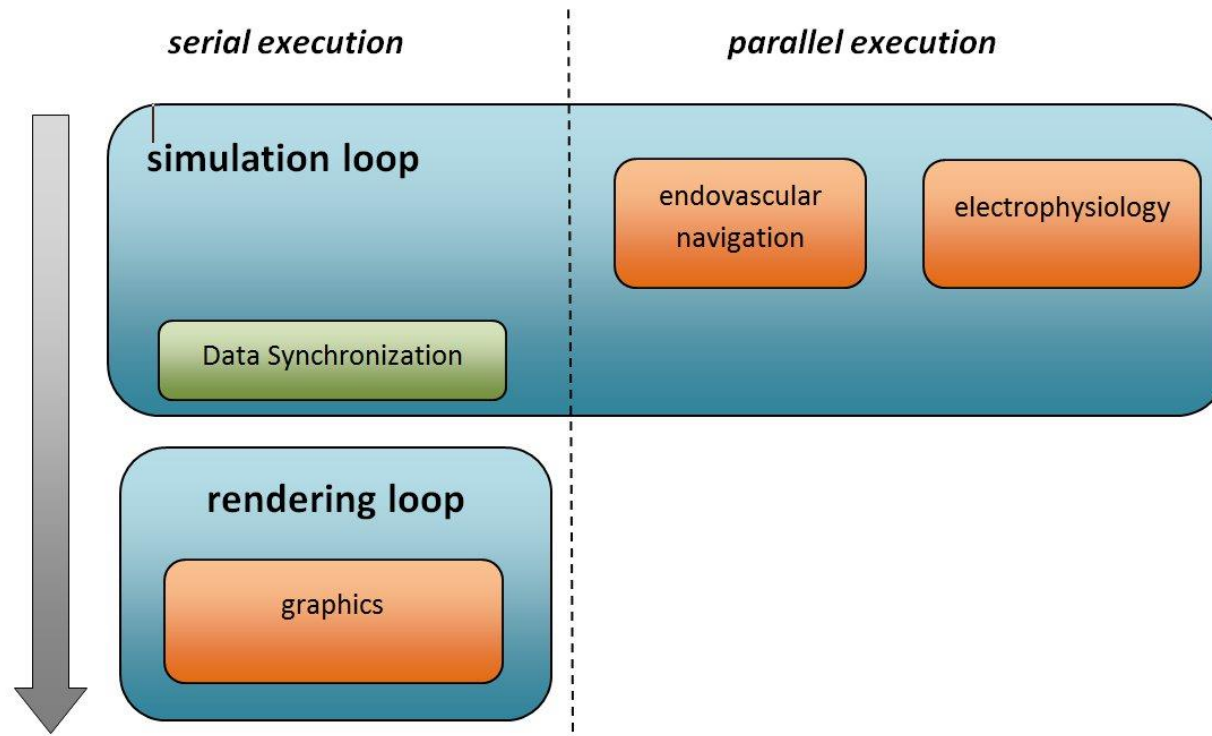


```
    <DataExchange source="@scenepath" dest="@scenepath" />  
  </node />
```


SOFA



Combining Multi-threading & GPU for RF Ablation Simulation




SOFA




AnimationLoopParallel Tips

each animation loop should be manually balanced to get almost the same duration:



replace the AnimationLoop of the fastest scenes with a *MultiStepAnimationLoop* and increase the number of *substep* untill to get a decrease of *fps*. Keep the time steps consistent




SOFA



AnimationLoopParallel Limitations

- No Collision interaction between objects in different scenes
- No mouse interaction with the objects in the scenes



BeamLinearMapping_mt (*for loop parallelization*)

- It inherits from *BeamLinearMapping* and overrides three virtual functions that contain a *for* loop: *apply()*, *applyJ()* and *applyJT()*
- Creates 3 task for each function overridden *applyTask*, *applyJTask* and *applyJTTask* and keeps locally in the task some shared data
- It adds the *granularity* attribute: the number of iterations of the *for* loop assigned and executed for each task.



Parallelize for loop

```
for ( i=0; i<N, ++i ) { some_code[i] }
```

- Define a *forTask* class with two indices members *first* and *last*
- Pass and initialize the *first* and *last* member in the constructor:

```
forTask::forTask( int first,int last, task::status* )
```

- Copy the loop in the run function using the *first* and *last* as range values:

```
forTask::run() { for ( i=first; i<last, ++i ) {some_code[i] } }
```


- Allocate the forTask(s) and add to the queue of the current thread

```
for ( i=0; i<N; i += grainSize )
```

```
currentThread->addTask( new forTask( i, i+grainsize, &status) )
```


- Wait for all the tasks completion

```
waitUntilDone(status);
```



Create a Parallel component

- Find the computationally intensive work with a performance profiler tool (vs analyzer, ..)
- decompose it into tasks, taking in account the parallelism level:
 - independent components that can run concurrently.
 - computationally intensive loop inside a function
- Check the share data and keep it local in the task.
- For big shared data arrays avoid or minimize the *false sharing* problem (keep more than 64bytes memory distance between the data acces from different cores)



Future work

- Add a *Parallel_For* function to parallelize automatically *for loop* tasks:

```
void Parallel_For( const forTask*, const GrainSize& )
```

- memory pool allocator for fast run-time tasks creation and destruction

